How can we define a *pretraining distribution* for DFAs?

Countably infinite concept class, any distribution is necessarily biased

Challenge: No Canonical Prior Distribution

Observation

Realizing any state transition through a DFA corresponds to a *reach-avoid problem*

Countably infinite concept class, any distribution is necessarily biased